



# StableHLO Compatibility

OpenXLA Community Meeting



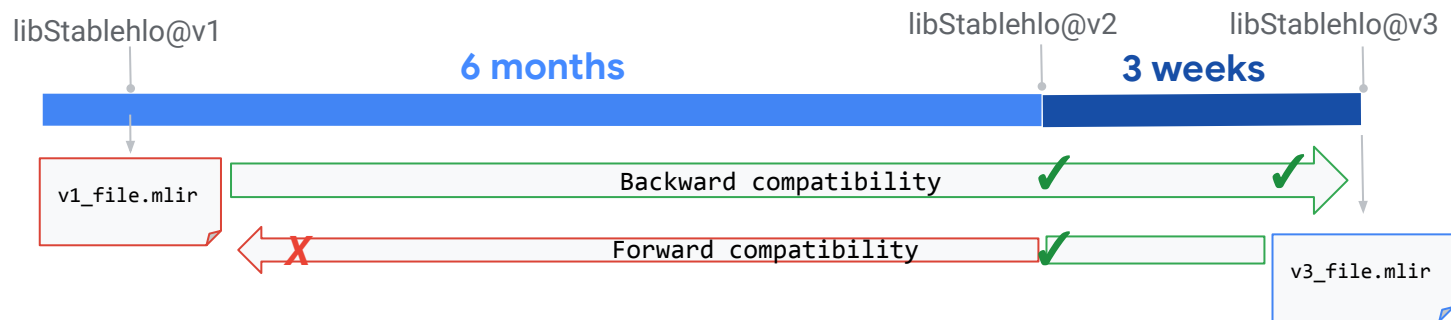
Kevin Gleason - 10/12/22

# Agenda

- Compatibility Requirements
- Upgrade and Downgrade Hooks
- Compatibility Prototype
- Next Steps

# Requirements (*feedback wanted!*)

R1	libStablehlo provides 6 months of backward compatibility
R2	libStablehlo provides 3 weeks of forward compatibility
R3	Source compatibility for C, C++ and Python APIs within libStablehlo is an aspirational goal.



# Supported Dialects and Operations, Attributes, Types

Dialect	Supported Ops, Attributes and Types
Arith	AddIOp, CmpIOp, CmpIPredicateAttr, ConstantOp, DivSIOp, ExtSIOp, IndexCastOp, MaxSIOp, MinSIOp, MulIOp, SelectOp, SubIOp, TruncIOp
Builtin	No ops, but all attributes and types.
CHLO	All ops, attributes and types.
Func	CallOp
MLProgram	All ops, attributes and types.
SparseTensor	Aspirational (pending the sparsity RFC which is expected in Q4 2022).
StableHLO	All ops, attributes and types except CustomCallOp whose semantics is implementation-defined.
Tensor	CastOp, DimOp, FromElementsOp

# IR Upgrades

- Provides backward compatibility
- For 6 months, support dialect upgrades
- Payloads have version number, allowing tooling to error if out of compatibility window

Rename `stablehlo.sub` → `stablehlo.subtract`

v39

```
func.func @upgrade(%arg0: tensor<2xi1>) -> () {  
  %0 = stablehlo.sub %arg0, %arg0 : tensor<2xi1>  
  return  
}
```

upgrade

v40

```
func.func @upgrade(%arg0: tensor<2xi1>) -> () {  
  %0 = stablehlo.subtract %arg0, %arg0 : tensor<2xi1>  
  return  
}
```

# IR Downgrades

- Provides forward compatibility when possible, else warn / error
- For 3 weeks after dialect changes, target the previous version
- Can be processed in v39, and can be upgraded by v40

Rename `stablehlo.sub` → `stablehlo.subtract`

v40

```
func.func @upgrade(%arg0: tensor<2xi1>) -> () {  
  %0 = stablehlo.subtract %arg0, %arg0 : tensor<2xi1>  
  return  
}
```

downgrade

v39

```
func.func @upgrade(%arg0: tensor<2xi1>) -> () {  
  %0 = stablehlo.sub %arg0, %arg0 : tensor<2xi1>  
  return  
}
```

# Compatibility Prototype Tool / API

## Compatibility Tool

```
stablehlo-translate --compat [flags] file.mlir
  --target=<version>    Target version to generate (Default to minimum supported)
  --emit-asm            Output textual assembly (Default generate bytecode)
```

## Compatibility APIs

```
// namespace mlir::stablehlo
OwningOpRef<Operation *> parseWithCompat(llvm::SourceMgr &sourceMgr,
                                          MLIRContext *context);

LogicalResult writeWithCompat(Operation *topLevelOperation,
                              int64_t targetVersion,
                              bool emitAssembly,
                              llvm::raw_ostream &output);
```

# Compatibility Prototype Example

```
void registerSubOpChanges() {
    // Change log:
    //   Version 39: SubOp<"stablehlo.sub"> exists
    //   Version 40: SubOp<"stablehlo.sub"> -> SubtractOp<"stablehlo.subtract">
    // Backward compatibility: Support v39 and after.
    // Forward compatibility: Target v39 for printing.

    // Upgrade <v40 -> 40: [sub --> subtract]
    addUpgrade("stablehlo.sub", 40,
        [&](Operation *op, int64_t fromVer) -> LogicalResult {
            renameOperation(op, "stablehlo.subtract");
            return success();
        });

    // Downgrade v40 -> 39: [subtract --> sub]
    addDowngrade("stablehlo.subtract", 39,
        [&](Operation *op, int64_t fromVer) -> LogicalResult {
            renameOperation(op, "stablehlo.sub");
            return success();
        });
}
```



# Compatibility Prototype Example

```
[gleasonk@gleasonk-linux:~/compat$ cat sub.mlir
func.func private @test_sub(%arg0: tensor<2xi1>) -> () attributes {stablehlo.compat_version = 37 : i32} {
  %0 = "stablehlo.sub"(%arg0, %arg0) : (tensor<2xi1>, tensor<2xi1>) -> tensor<2xi1>
  func.return
}

[gleasonk@gleasonk-linux:~/compat$ stablehlo-translate --compat --target=40 sub.mlir
func.func private @test_sub(%arg0: tensor<2xi1>) attributes {stablehlo.compat_version = 40 : i64} {
  %0 = "stablehlo.subtract"(%arg0, %arg0 : tensor<2xi1>)
  return
}

[gleasonk@gleasonk-linux:~/compat$ stablehlo-translate --compat --target=38 sub.mlir
"func.func"() ({
^bb0(%arg0: tensor<2xi1>):
  %0 = "stablehlo.sub"(%arg0, %arg0) : (tensor<2xi1>, tensor<2xi1>) -> tensor<2xi1>
  "func.return"() : () -> ()
}) {function_type = (tensor<2xi1>) -> (), stablehlo.compat_version = 38 : i64, sym_name = "test_sub", sym_
) -> ()
```

# Next Steps

- In Q4 2022:
  - Propose an MLIR RFC about stability of the bytecode format.
  - Start a discussion about stability guarantees for the Builtin dialect.
  - Finish the compatibility tool, hold a design review for the StableHLO Compatibility RFC
- We aim to ship a working solution by the end of the year.